

Технології машинного навчання, лаб 1

Пашковський Євгеній, ІІ-41мн

**task1** code:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
import tensorflow as tf

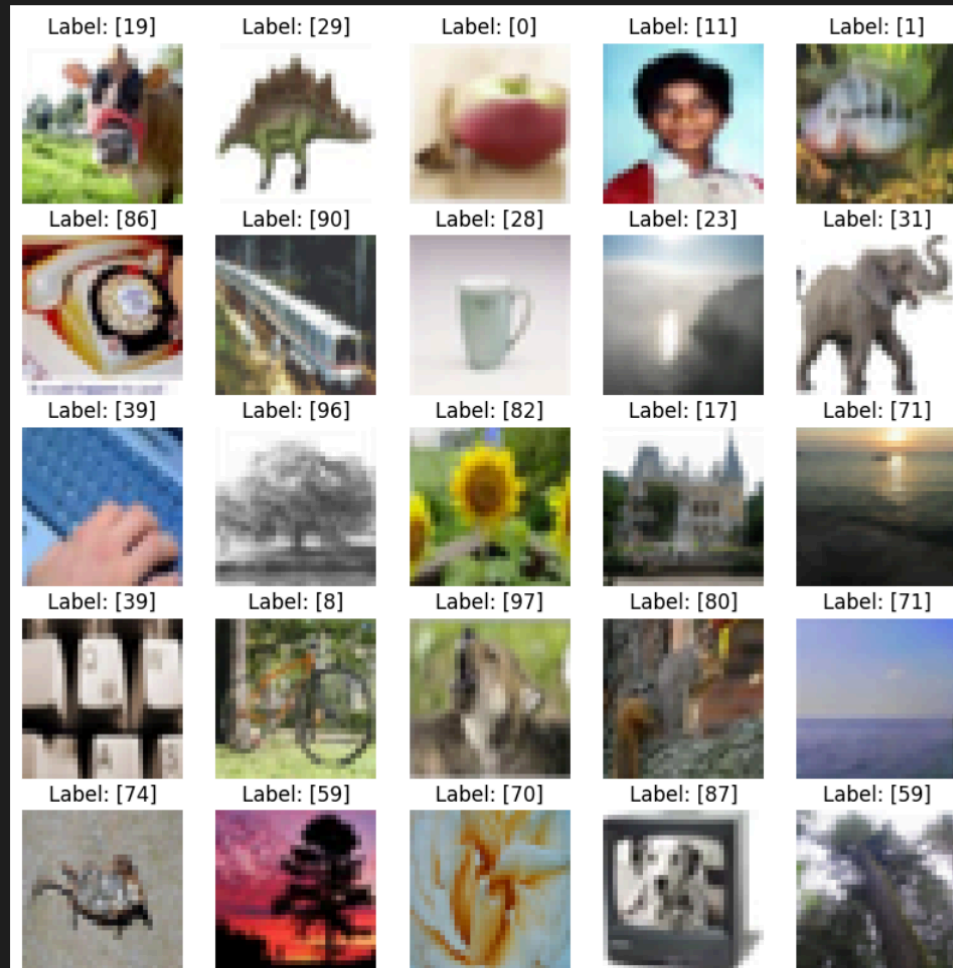
(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.cifar100.load_data()

print(f'x_train shape: {x_train.shape}')
print(f'x_test shape: {x_test.shape}')

print(f'y_train shape: {y_train.shape}')
print(f'y_test shape: {y_test.shape}')

plt.figure(figsize=(10, 10))
for i in range(25):
    ax = plt.subplot(5, 5, i + 1)
    plt.imshow(x_train[i])
    plt.title(f'Label: {y_train[i]}')
    plt.axis('off')
plt.show()
```

```
x_train shape: (50000, 32, 32, 3)
x_test shape: (10000, 32, 32, 3)
y_train shape: (50000, 1)
y_test shape: (10000, 1)
```



```
import keras
from keras import models, layers

def MyConv2D(x, filters = 32, kernel_size = (3, 3), strides = (1, 1),
padding = "same", activation = "relu"):
    x = layers.Conv2D(
        filters=filters,
        kernel_size=kernel_size,
        strides=strides,
        padding=padding,
    )(x)

    x = layers.BatchNormalization(axis=-1)(x)
    x = layers.Activation("relu")(x)
    return x
```

```

def Stem(input):
    x = MyConv2D(input, filters=32, kernel_size=(3,3), strides=(2,2),
padding="valid", activation="relu")
    x = MyConv2D(x, filters=32, kernel_size=(3,3), padding="valid",
activation="relu")
    x = MyConv2D(x, filters=64, kernel_size=(3,3), padding="same",
activation="relu")
    x = layers.MaxPooling2D((3, 3), strides=(2, 2), padding="valid")(x)
    x = MyConv2D(x, filters=80, kernel_size=(1,1), padding="valid",
activation="relu")
    x = MyConv2D(x, filters=192, kernel_size=(3,3), padding="valid",
activation="relu")
    x = layers.MaxPooling2D((3, 3), strides=(2, 2), padding="valid")(x)
    return x

def InceptionA(input, filters=64):
    x1 = MyConv2D(input, filters=64, kernel_size=(1,1), padding="same",
activation="relu")
    x1 = MyConv2D(x1, filters=96, kernel_size=(3,3), padding="same",
activation="relu")
    x1 = MyConv2D(x1, filters=96, kernel_size=(3,3), padding="same",
activation="relu")

    x2 = MyConv2D(input, filters=48, kernel_size=(1,1), padding="same",
activation="relu")
    x2 = MyConv2D(x2, filters=64, kernel_size=(3,3), padding="same",
activation="relu")

    x3 = layers.AveragePooling2D((3, 3), strides=(1, 1),
padding="same")(input)
    x3 = MyConv2D(x3, filters=filters, kernel_size=(1,1), padding="same",
activation="relu")

    x4 = MyConv2D(input, filters=64, kernel_size=(1,1), padding="same",
activation="relu")

    x = layers.Concatenate(axis=-1)([x1, x2, x3, x4])
    return x

def ReductionA(input):
    x1 = MyConv2D(input, filters=64, kernel_size=(1,1), padding="same",
activation="relu")

```

```

        x1 = MyConv2D(x1, filters=96, kernel_size=(3,3), padding="same",
activation="relu")
        x1 = MyConv2D(x1, filters=96, kernel_size=(3,3), strides=(2,2),
padding="valid", activation="relu")

        x2 = MyConv2D(input, filters=384, kernel_size=(3,3), strides=(2,2),
padding="valid", activation="relu")

        x3 = layers.MaxPooling2D((3, 3), strides=(2, 2),
padding="valid")(input)

        x = layers.Concatenate(axis=-1)([x1, x2, x3])
        return x

def InceptionB(input, filters=192):
    x1 = MyConv2D(input, filters=filters, kernel_size=(1,1),
padding="same", activation="relu")
    x1 = MyConv2D(x1, filters=filters, kernel_size=(7,1), padding="same",
activation="relu")
    x1 = MyConv2D(x1, filters=filters, kernel_size=(1,7), padding="same",
activation="relu")
    x1 = MyConv2D(x1, filters=filters, kernel_size=(7,1), padding="same",
activation="relu")
    x1 = MyConv2D(x1, filters=192, kernel_size=(1,7), padding="same",
activation="relu")

    x2 = MyConv2D(input, filters=filters, kernel_size=(1,1),
padding="same", activation="relu")
    x2 = MyConv2D(x2, filters=filters, kernel_size=(1,7), padding="same",
activation="relu")
    x2 = MyConv2D(x2, filters=192, kernel_size=(7,1), padding="same",
activation="relu")

    x3 = layers.AveragePooling2D((3, 3), strides=(1, 1),
padding="same")(input)
    x3 = MyConv2D(x3, filters=192, kernel_size=(1,1), padding="same",
activation="relu")

    x4 = MyConv2D(input, filters=192, kernel_size=(1,1), padding="same",
activation="relu")

    x = layers.Concatenate(axis=-1)([x1, x2, x3, x4])

```

```

        return x

    def ReductionB(input):
        x1 = MyConv2D(input, filters=192, kernel_size=(1,1), padding="same",
activation="relu")
        x1 = MyConv2D(x1, filters=192, kernel_size=(1,7), padding="same",
activation="relu")
        x1 = MyConv2D(x1, filters=192, kernel_size=(7,1), padding="same",
activation="relu")
        x1 = MyConv2D(x1, filters=192, kernel_size=(3,3), strides=(2,2),
padding='valid', activation="relu")

        x2 = MyConv2D(input, filters=192, kernel_size=(1,1), padding="same",
activation="relu")
        x2 = MyConv2D(input, filters=320, kernel_size=(3,3), strides=(2,2),
padding='valid', activation="relu")

        x3 = layers.MaxPooling2D((3, 3), strides=(2, 2),
padding="valid")(input)

        x = layers.Concatenate(axis=-1)([x1, x2, x3])
        return x

    def InceptionC(input):
        x1 = MyConv2D(input, filters=448, kernel_size=(1,1), padding="same",
activation="relu")
        x1 = MyConv2D(x1, filters=384, kernel_size=(3,3), padding="same",
activation="relu")
        x1_1 = MyConv2D(x1, filters=384, kernel_size=(1,3), padding="same",
activation="relu")
        x1_2 = MyConv2D(x1, filters=384, kernel_size=(3,1), padding="same",
activation="relu")

        x2 = MyConv2D(input, filters=384, kernel_size=(1,1), padding="same",
activation="relu")
        x2_1 = MyConv2D(x2, filters=384, kernel_size=(1,3), padding="same",
activation="relu")
        x2_2 = MyConv2D(x2, filters=384, kernel_size=(3,1), padding="same",
activation="relu")

```

```

        x3 = layers.AveragePooling2D((3, 3), strides=(1, 1),
padding="same")(input)
        x3 = MyConv2D(x3, filters=192, kernel_size=(1,1), padding="same",
activation="relu")

        x4 = MyConv2D(input, filters=320, kernel_size=(1,1), padding="same",
activation="relu")

        x = layers.Concatenate(axis=-1)([x1_1, x1_2, x2_1, x2_2, x3, x4])
        return x

def Aux(input, n_classes = 1000):
    x = layers.AveragePooling2D((5, 5), strides=(3, 3),
padding="same")(input)
    x = MyConv2D(x, filters=128, kernel_size=(1,1), padding="same",
activation="relu")
    x = layers.Flatten()(x)
    x = layers.Dense(768, activation="relu")(x)
    x = layers.Dropout(rate = 0.2)(x)
    x = layers.Dense(n_classes, activation="softmax")(x)
    return x

def InceptionV3(input_shape=(299, 299, 3), n_classes = 1000):
    input = layers.Input(input_shape)
    x = layers.Rescaling(1./255)(input)
    x = layers.Resizing(299, 299)(x)
    x = layers.RandomColorJitter(
        value_range=(0, 1),
        brightness_factor=0.2,
        contrast_factor=0.2,
        saturation_factor=0.2,
        hue_factor=0.2
    )(x)
    x = layers.RandomFlip(
        mode="horizontal_and_vertical"
    )(x)

    x = Stem(x)

    x = InceptionA(x, 32)
    x = InceptionA(x)
    x = InceptionA(x)

```

```

x = ReductionA(x)

x = InceptionB(x, 128)
x = InceptionB(x, 160)
x = InceptionB(x, 160)
x = InceptionB(x, 192)

aux = Aux(x, n_classes)

x = ReductionB(x)

x = InceptionC(x)
x = InceptionC(x)

x = layers.GlobalAveragePooling2D()(x)

x = layers.Dense(2048, activation="relu")(x)
x = layers.Dropout(rate = 0.2) (x)
output = layers.Dense(n_classes, activation="softmax")(x)

model = models.Model(inputs=input, outputs=output)
return model

```

```

from keras import optimizers

epochs = 50

initial_learning_rate = 1e-3
final_learning_rate = 1e-8
learning_rate_decay_factor = (final_learning_rate /
initial_learning_rate) ** (1 / epochs)
batch_size = 32
steps_per_epoch = len(x_train) * 8 // 10 // batch_size

print(f'Initial learning rate: {initial_learning_rate}')
print(f'Final learning rate: {final_learning_rate}')
print(f'Learning rate decay factor: {learning_rate_decay_factor}')
print(f'Batch size: {batch_size}')
print(f'Steps per epoch: {steps_per_epoch}')

```

```

learning_rate = optimizers.schedules.ExponentialDecay(
    initial_learning_rate=initial_learning_rate,
    decay_steps=steps_per_epoch,
    decay_rate=learning_rate_decay_factor
)

model = InceptionV3(input_shape=(32, 32, 3), n_classes=100)
model.compile(optimizer=optimizers.Adam(learning_rate=learning_rate),
loss='sparse_categorical_crossentropy', metrics=['accuracy'])

model.summary()

```

```

history = model.fit(x_train, y_train, epochs=epochs,
validation_split=0.2, steps_per_epoch=steps_per_epoch, batch_size=batch_size)

```

```

model.evaluate(x_test, y_test)

```

```

313/313 ————— 11s 34ms/step - accuracy: 0.5111 - loss: 1.9735
[2.0017733573913574, 0.5048999786376953]

```

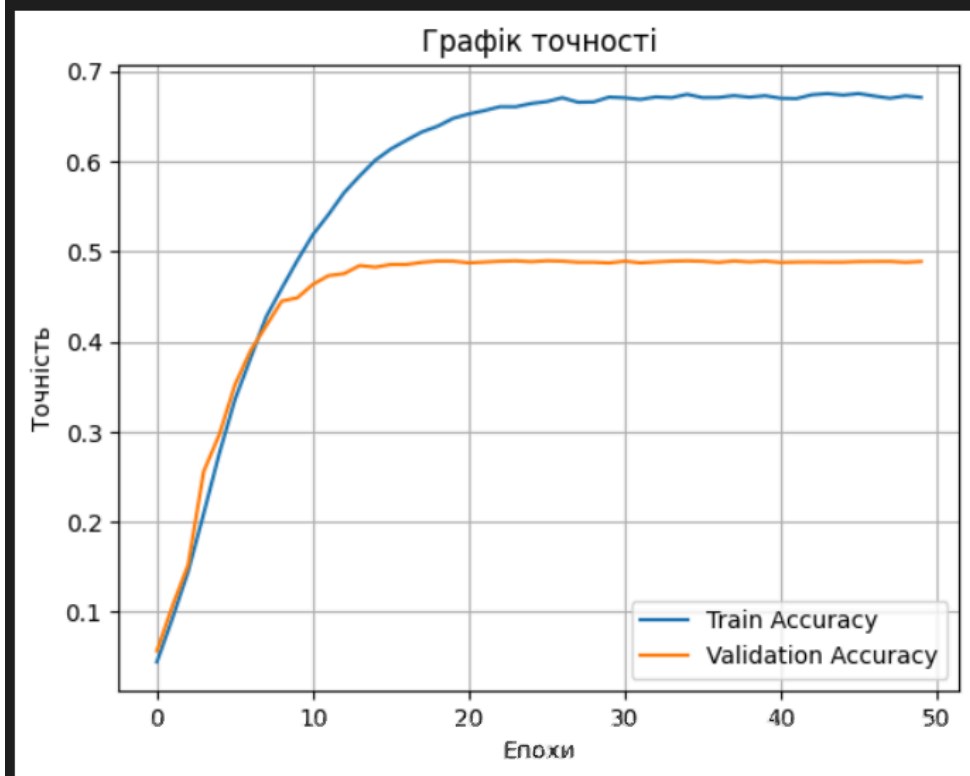
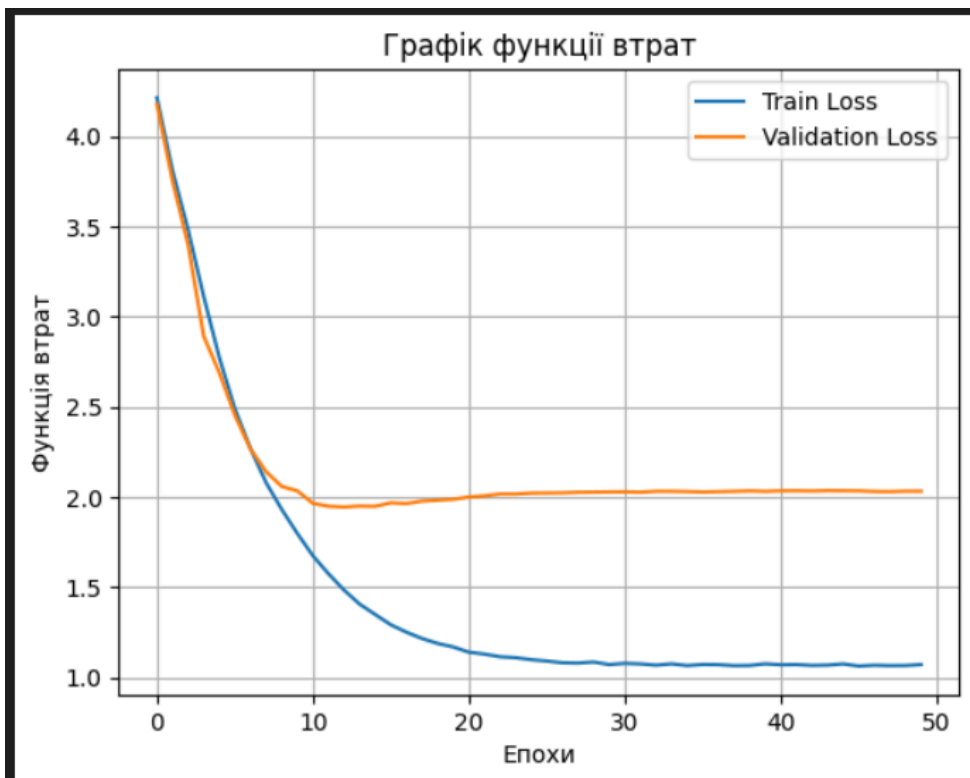
```

plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Епохи')
plt.ylabel('Функція втрат')
plt.title('Графік функції втрат')
plt.legend()
plt.grid()
plt.show()

plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Епохи')
plt.ylabel('Точність')
plt.title('Графік точності')
plt.legend()
plt.grid()
plt.show()

```





```
predictions = model.predict(x_test).argmax(axis=1)
```

```
predictions
```

```
models.save_model(model, 'model.keras')
```

```
model = models.load_model('model.keras')  
model.evaluate(x_test, y_test)
```

**task2** code for .tflite file generation:

```
from keras import models  
import tensorflow as tf  
  
model = models.load_model('../task1/model.keras')
```

```
converter = tf.lite.TFLiteConverter.from_keras_model(model)  
tflite_model = converter.convert()  
  
with tf.io.gfile.GFile('./model/model.tflite', 'wb') as f:  
    f.write(tflite_model)
```

**task2** code for server with /predict endpoint:

```
from fastapi import FastAPI, File, UploadFile  
from fastapi.responses import JSONResponse  
from PIL import Image  
import io  
import numpy as np  
import tensorflow as tf  
  
app = FastAPI()  
  
cifar100_fine_labels = [  
    'apple', 'aquarium_fish', 'baby', 'bear', 'beaver', 'bed', 'bee',  
'beetle',  
    'bicycle', 'bottle', 'bowl', 'boy', 'bridge', 'bus', 'butterfly',  
    'camel', 'can', 'castle', 'caterpillar', 'cattle', 'chair',  
'chimpanzee',  
    'clock', 'cloud', 'cockroach', 'couch', 'crab', 'crocodile', 'cup',  
'dinosaur',  
    'dolphin', 'elephant', 'flatfish', 'forest', 'fox', 'girl',  
'hamster', 'house',
```

```

        'kangaroo', 'keyboard', 'lamp', 'lawn_mower', 'leopard', 'lion',
        'lizard',
        'lobster', 'man', 'maple_tree', 'motorcycle', 'mountain', 'mouse',
        'mushroom',
        'oak_tree', 'orange', 'orchid', 'otter', 'palm_tree', 'pear',
        'pickup_truck',
        'pine_tree', 'plain', 'plate', 'poppy', 'porcupine', 'possum',
        'rabbit',
        'raccoon', 'ray', 'road', 'rocket', 'rose', 'sea', 'seal', 'shark',
        'shrew',
        'skunk', 'skyscraper', 'snail', 'snake', 'spider', 'squirrel',
        'streetcar',
        'sunflower', 'sweet_pepper', 'table', 'tank', 'telephone',
        'television', 'tiger',
        'tractor', 'train', 'trout', 'tulip', 'turtle', 'wardrobe', 'whale',
        'willow_tree',
        'wolf', 'woman', 'worm'
    ]

    print(len(cifar100_fine_labels))

    model_content = open("./model/model.tflite", "rb").read()
    interpreter = tf.lite.Interpreter(model_content=model_content)

    def infer(img):
        interpreter.allocate_tensors()

        input_details = interpreter.get_input_details()
        output_details = interpreter.get_output_details()

        input_data = np.array([np.array(img)], dtype=np.float32)
        interpreter.set_tensor(input_details[0]['index'], input_data)

        interpreter.invoke()

        output_data = interpreter.get_tensor(output_details[0]['index'])
        return output_data

    @app.get("/ping")
    async def ping():

        return "pong"

```

```

@app.post("/predict")
async def predict(file: UploadFile = File(...)):
    if not file.content_type.startswith("image/"):
        return JSONResponse(
            status_code=400,
            content={"error": "Invalid file type. Only image files are
allowed."}
        )

    contents = await file.read()
    image = Image.open(io.BytesIO(contents))

    predictions = infer(image)
    labelIndex = int(np.array(predictions).argmax())

    return {
        "prediction": labelIndex,
        "label": cifar100_fine_labels[labelIndex]
    }

```

## task2 Dockerfile:

```

FROM tensorflow/tensorflow
WORKDIR /app
COPY . .
RUN pip install --upgrade pip
RUN pip install fastapi pillow uvicorn numpy python-multipart
tensorflow-cpu
EXPOSE 3000
CMD ["uvicorn", "--host", "0.0.0.0", "--port", "3000", "server:app"]

```

## task2 process of running and querying the api:

1. build docker image

```

PS E:\Marictp\semestr2\machine-learning-technologies> cd .\L1\task2\
PS E:\Marictp\semestr2\machine-learning-technologies\L1\task2> docker build -t l1task2:1.0 .
[+] Building 56.5s (10/10) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 283B
=> [internal] load metadata for docker.io/tensorflow/tensorflow:latest
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/5] FROM docker.io/tensorflow/tensorflow:latest@sha256:f24e8494d43a4e613edd7fbd3782594368d52d79af1e216051139ed2d7830682
=> [internal] load build context
=> => transferring context: 626.51kB
=> CACHED [2/5] WORKDIR /app
=> [3/5] COPY . .
=> [4/5] RUN pip install --upgrade pip
=> [5/5] RUN pip install fastapi pillow uvicorn numpy python-multipart tensorflow-cpu
=> exporting to image
=> => exporting layers
=> => writing image sha256:4db0774acf8a5768be68eb57fba9e938c07c6e2562e20297199f5cd5c13b0f81
=> => naming to docker.io/library/l1task2:1.0

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/7qe7sia8un2obuusi9qamhxy0

What's next:
  View a summary of image vulnerabilities and recommendations → docker scout quickview
PS E:\Marictp\semestr2\machine-learning-technologies\L1\task2>

```

## 2. run docker container & expose port

```

PS E:\Marictp\semestr2\machine-learning-technologies\L1\task2> docker run -p 3000:3000 l1task2:1.0
2025-05-16 03:17:13.971795: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You
. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-05-16 03:17:16.002878: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow bin
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow wit
2025-05-16 03:17:34.548278: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You
. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
/usr/local/lib/python3.11/dist-packages/tensorflow/lite/python/interpreter.py:457: UserWarning:
  TF 2.20. Please use the LiteRT interpreter from the ai_edge_litert package.
  See the [migration guide](https://ai.google.dev/edge/litert/migration)
  for details.

warnings.warn(_INTERPRETER_DELETION_WARNING)
INFO:      Started server process [1]
INFO:      Waiting for application startup.
INFO:      Application startup complete.
INFO:      Uvicorn running on http://0.0.0.0:3000 (Press CTRL+C to quit)


```


## 3. Enter url & attach image file in postman

The screenshot shows the Postman interface for a POST request to `localhost:3000/predict`. The 'Body' tab is selected, and the 'form-data' radio button is chosen. A single form data entry is visible:

| Key  | Value                    | Description |
|------|--------------------------|-------------|
| file | dwarf_tulip_s_000032.png |             |

Below this table, there is a section for additional key-value pairs with 'Key' and 'Value' headers.

ary  GraphQL

|      | Value  |
|------|--|
| e ▾  |  dwarf_tulip_s_000032.png |
| xt ▾ | Value  |

4. Make request POST /predict with image attached

```
1 {  
2   "prediction": 92,  
3   "label": "tulip"  
4 }
```

Result is correct!